

Yet Another Great Ontology (YAGO)

by Ole Berg & Lennard Feuerbach
on the 5th of November, 2024

Table of Contents

1. Introduction
2. Versions of YAGO
3. Applications
4. YAGO and SPARQL
5. Comparison
6. Conclusion

1. Introduction

1.1 Ontologies (Recap):

- Formal Representation: Structured definitions of concepts and relationships
- Shared Vocabulary: Common language for effective data exchange
- Knowledge Integration: Unifies diverse data sources for a cohesive view
- Reasoning and Inference: Enables logical reasoning to derive new knowledge

1. Introduction

1.2 Motivation and history

- Researchers at the Max Planck Institute for Informatics in Saarbrücken saw potential of knowledge bases
- Already had various application fields
- Criticized that applications only utilized a single source of knowledge, often either WordNet or Wikipedia
- Laid out what features and traits a heavily improved application should have:
 - High accuracy (close to 100 %)
 - Combination of concepts and entities
 - Extensible, easily re-usable, and application independent

1. Introduction

1.3 Yet Another Great Ontology (YAGO):

- Knowledge Base: Contains well over millions of entities and facts from real world
- Entities and Relations: Includes entities and their relationships
 - e.g. who played in which movie, which city is located in which country
- Class Taxonomy: Organizes entities in different hierarchical classes
 - e.g. class of cities is a subclass of the class of populated places
- Defined Relations: Specifies relations between entities, forming the ontology
 - e.g. birthPlace can only be between person and a place

1. Introduction

1.3 Yet Another Great Ontology (YAGO):

- Wikidata's vast facts are combined with schema.org's ontology
- All identifiers are human-readable and comparable easy to understand
- Logical constraints maintain data quality and enables reasoning over the data
- Simpler version of Wikidata, offering a cleaner knowledge base

2. Versions of YAGO

2.1 YAGO 1

- First version of YAGO, created in 2007
- Unifies facts from Wikipedia with clean taxonomy from WordNet
- Categories from Wikipedia are brought together with nouns from WordNet
- 2 million entities and 20 million facts

2. Versions of YAGO

2.2 YAGO2

- Improvement in 2012
- “anchoring [...] along both the spatial and the temporal dimensions” (Hoffart et al. 2012, p. 3)
- *When and where was a fact true?*
- Introduction of *SPOTL(X)* for representing and querying “spatio-temporally enhanced facts” (Hoffart et al. 2012, p. 3)
- Time (or time frame) for entities and facts
- Single geo-coordinate pair for *geo-entities*
- 9.8 million entities and 447 million facts
- Coverage: Time 47 %, location 30 %

2. Versions of YAGO

2.3 YAGO2s

- Refactoring of YAGO2 in 2013
- Transparent and modular architecture
- Compliance of YAGO syntax with RDF
- Usage of Turtle format
- Introduction of domains from WordNet

2. Versions of YAGO

2.4 YAGO3

- Improvement in 2015
- Expansion into Wikipedias in other languages
- Use of Wikidata to identify same entities in different-language Wikipedias
- Fact extraction from infoboxes in Wikipedia
- 17 million entities and more than 150 million facts

2. Versions of YAGO

2.5 YAGO4

- Complete overhaul in 2020
- WordNet is abandoned, replaced by Schema.org
- Focus on Wikidata instead of Wikipedia infoboxes
- Browser and SPARQL endpoint are available
- Upper classes from Schema.org and selected lower classes from Wikidata
- 67 million entities and 343 million facts (Full YAGO)

2. Versions of YAGO

2.6 YAGO4.5

- Addresses issues of YAGO4
- Paper has been published this year (2024)
- Important principles:
 - Upper taxonomy to define properties
 - Lower taxonomy to convey information in human-readable form
- 49 million entities and 132 million facts

3. Applications

3.1 Various examples

- NER: Enhances entity recognition with structured data
- Semantic Web: Builds knowledge graphs for data linking
- AI Chatbots: Improves chatbot responses with factual data
- Knowledge Graphs: Maps complex entity relationships
- NLP: Supports semantic analysis and word disambiguation

3. Applications

3.2 IBM Watson

- The most known application of YAGO is IBM Watson
- IBM Watson is a system capable of answering questions
- Developed by at IBM, team led by David Ferucci, released in 2011
- Designed to answer questions from the TV show Jeopardy
- IBM Watson actually beat the real champions and won

3. Applications

3.2 IBM Watson

- Created as a question-answering system using NLP, IR and other technologies
- Written in languages including Java, C++ and Prolog, and runs on Linux servers
- The system Watson uses has 2,880 processor threads and 16 TB of RAM
- The data Watson uses includes encyclopaedias, dictionaries, articles and more
- Also databases, taxonomies and ontologies as DBPedia, WordNet and **YAGO**

4. YAGO and SPARQL

```
PREFIX schema: <http://schema.org/>
```

```
PREFIX yago: <http://yago-knowledge.org/resource/>
```

```
PREFIX rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#>
```

```
PREFIX xsd: <http://www.w3.org/2001/XMLSchema#>
```

```
SELECT ?movie ?duration WHERE {  
  ?movie rdf:type schema:Movie .  
  ?movie schema:duration ?duration .  
  FILTER (xsd:decimal(?duration) >= 100)  
}  
LIMIT 3
```


4. YAGO and SPARQL

	movie		duration
1	yago:The_Way_of_the_Dragon		"100"^^xsd:decimal
2	yago:Carrie__u0028_2013_film_u0029_		"100"^^xsd:decimal
3	yago:The_Profession_of_Arms__u0028_2001_film_u0029_		"100"^^xsd:decimal

Figure 1: Result set of the SPARQL query

Try it for yourself: <https://yago-knowledge.org/sparql>

5. Comparison

- YAGO vs. DBpedia: Fixed schema, non-redundant relations, 50m vs. 4m instances
- YAGO vs. ConceptNet: Focus on specific instances, not common sense knowledge
- YAGO vs. BabelNet: Structured taxonomy and defined schema
- YAGO vs. Freebase: Actively maintained with ongoing updates
- YAGO vs. Wikidata: Readable identifiers, clean taxonomy, logical constraints

6. Conclusion

- **Combines Strengths:** Merges Wikidata's facts with schema.org's ontology
- **User-Friendly:** Has human-readable identifiers and well-organized classifications
- **Streamlined Data:** Includes only well-populated, essential classes and properties
- **Logical Constraints:** Ensures data quality and supports advanced reasoning

7. References

- F. M. Suchanek, G. Kasneci, und G. Weikum, “YAGO: A Core of Semantic Knowledge Unifying WordNet and Wikipedia”, 2007.
- J. Hoffart, F. M. Suchanek, K. Berberich, und G. Weikum, “YAGO2: A spatially and temporally enhanced knowledge base from Wikipedia”, Artificial Intelligence, Bd. 194, S. 28-61, Jan. 2013, doi: 10.1016/j.artint.2012.06.001.
- J. Biega, E. Kuzey, und F. M. Suchanek, “Inside YAGO2s: a transparent information extraction architecture”, in Proceedings of the 22nd International Conference on World Wide Web, Rio de Janeiro Brazil: ACM, Mai 2013, S. 325-328. doi: 10.1145/2487788.2487935.
- F. Mahdisoltani, J. Biega, und F. M. Suchanek, “YAGO3: A Knowledge Base from Multilingual Wikipedias”.

7. References (cont.)

- T. Pellissier Tanon, G. Weikum, und F. Suchanek, “YAGO 4: A Reason-able Knowledge Base”, in The Semantic Web, Bd. 12123, A. Harth, S. Kirrane, A.-C. Ngonga Ngomo, H. Paulheim, A. Rula, A. L. Gentile, P. Haase, und M. Cochez, Hrsg., in Lecture Notes in Computer Science, vol. 12123. , Cham: Springer International Publishing, 2020, S. 583-596. doi: 10.1007/978-3-030-49461-2_34.
- F. M. Suchanek, M. Alam, T. Bonald, L. Chen, P.-H. Paris, und J. Soria, “YAGO 4.5: A Large and Clean Knowledge Base with a Rich Taxonomy”, in Proceedings of the 47th International ACM SIGIR Conference on Research and Development in Information Retrieval, Washington DC USA: ACM, Juli 2024, S. 131-140. doi: 10.1145/3626772.3657876.

7. References (cont.)

- Ferrucci, David & Brown, Eric & Chu-Carroll, Jennifer & Fan, James & Gondek, David & Kalyanpur, Aditya & Lally, Adam & Murdock, J William & Nyberg, Eric & Prager, John & Schlaefher, Nico & Welty, Christopher. (2010). Building Watson: An Overview of the DeepQA Project. AI Magazine. 31. 59-79.
10.1609/aimag.v31i3.2303.
- “YAGO 3.0”. Zugriff am 01.11.2024 [Online]. Link:
<https://github.com/yago-naga/yago3>
- “YAGO 4.5”. Zugriff am 01.11.2024 [Online]. Link:
<https://github.com/yago-naga/yago-4.5>
- “Getting Started”. Zugriff am 01.11.2024 [Online]. Link:
<https://yago-knowledge.org/getting-started>